

MATH 345 Skeleton Notes

Unit 1: Vectors, matrices, and linear maps

Section 1.3: Matrix-vector product and linear maps

© 2024–2026 Jacob Denson, Kaitlyn Phillipson, Sebastien Roch. Except where otherwise noted, this work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

Matrix-vector product

A vector can be written as a row or as a column. When a matrix acts on a vector in this course, we usually write the vector as a column:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}.$$

From now on, when we write $\mathbf{x} = (x_1, \dots, x_n)$, $\mathbf{x}$ is interpreted as a column vector.

If A is an $m \times n$ matrix, then A has n columns. We write

$$A = [\mathbf{a}_1 \ \mathbf{a}_2 \ \cdots \ \mathbf{a}_n]$$

where each $\mathbf{a}_j$ is a column vector in $\mathbb{R}^m$.

Definition: Matrix-vector product

Let $A = [\mathbf{a}_1 \ \mathbf{a}_2 \ \cdots \ \mathbf{a}_n]$ be an $m \times n$ matrix and let $\mathbf{x} = (x_1, \dots, x_n)$ be a column vector in $\mathbb{R}^n$. The is

$$A\mathbf{x} = x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n.$$

Activity: Computing a matrix-vector product (U1-LO4)

Compute $A\mathbf{x}$ where

$$A = \begin{bmatrix} -1 & 4 & -5 \\ 3 & 1 & -2 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 2 \\ -3 \\ 4 \end{bmatrix}.$$

Theorem: Row dot-product view

If A is an $m \times n$ matrix with rows $\text{row}_1(A), \dots, \text{row}_m(A)$ as n -vectors, then

$$A\mathbf{x} = \begin{bmatrix} \text{row}_1(A) \cdot \mathbf{x} \\ \text{row}_2(A) \cdot \mathbf{x} \\ \vdots \\ \text{row}_m(A) \cdot \mathbf{x} \end{bmatrix}.$$

Why is this true?

Write $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$, where $\mathbf{a}_j = (a_{1j}, \dots, a_{mj})$, and let $\mathbf{x} = (x_1, \dots, x_n)$. Since

$$A\mathbf{x} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n,$$

its i -th entry is

$$x_1a_{i1} + \cdots + x_na_{in} = \text{row}_i(A) \cdot \mathbf{x}.$$

Stacking these entries gives the stated formula.

Activity: Computing the same product using rows (U1-LO4)

We revisit Computing a matrix-vector product (U1-LO4) using row dot products. Compute $A\mathbf{x}$ where

$$A = \begin{bmatrix} -1 & 4 & -5 \\ 3 & 1 & -2 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 2 \\ -3 \\ 4 \end{bmatrix}.$$

Theorem: Properties of matrix-vector products

Let A and B be $m \times n$ matrices, let $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$, and let r be a scalar. Then

- $A(\mathbf{x} + \mathbf{y}) = A\mathbf{x} + A\mathbf{y}$,
- $A(r\mathbf{x}) = r(A\mathbf{x}) = (rA)\mathbf{x}$,
- $(A + B)\mathbf{x} = A\mathbf{x} + B\mathbf{x}$.

Why is this true?

We prove the first property. Write $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$, $\mathbf{x} = (x_1, \dots, x_n)$, and $\mathbf{y} = (y_1, \dots, y_n)$. By the definition above,

$$\begin{aligned} A(\mathbf{x} + \mathbf{y}) &= (x_1 + y_1)\mathbf{a}_1 + \cdots + (x_n + y_n)\mathbf{a}_n \\ &= A\mathbf{x} + A\mathbf{y}. \end{aligned}$$

Note: Shape habit for matrix-vector products

Before forming $A\mathbf{x}$, check that the number of columns of A equals the number of entries of $\mathbf{x}$. If A is $m \times n$ and $\mathbf{x} \in \mathbb{R}^n$, then $A\mathbf{x} \in \mathbb{R}^m$: its m entries come from the m row dot products above.

Note: Rows measure; columns contribute

A matrix-vector product has two complementary readings. In the row view, each row measures the input by taking a dot product with $\mathbf{x}$. In the column view,

$$A\mathbf{x} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n,$$

so the input coordinates tell how much each column contributes to the output.

Examples of matrix-vector products

The same product Ax can represent different operations, depending on what the rows and columns of A mean. The next activities show three common patterns: scoring, selecting, and measuring changes.

Activity: Scoring objects by features (U1-LO4)

Let the rows of X represent three objects with two features:

$$X = \begin{bmatrix} 3 & 1 \\ 1 & 4 \\ 2 & 2 \end{bmatrix}, \quad \mathbf{w} = (2, -1).$$

Compute $X\mathbf{w}$. Interpret the result.

Activity: A selector matrix (U1-LO4, U1-LO5)

Let

$$S = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 5 \\ 6 \\ 7 \\ 8 \end{bmatrix}.$$

Compute $S\mathbf{x}$. What does S do?

Activity: A difference matrix (U1-LO4, U1-LO5)

Let

$$D = \begin{bmatrix} -1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 2 \\ 5 \\ 9 \end{bmatrix}.$$

Compute $D\mathbf{x}$. What does D measure?

Activity: Document ranking in matrix code (U1-LO4, U1-LO8)

In Word-count vectors (U1-LO1, U1-LO2), we ranked documents by cosine similarity one score at a time. The following code repeats that ranking using a data matrix whose rows are the document vectors.

```
import numpy as np

doc_names = np.array(["D1", "D2", "D3"])

X = np.array([
    [3., 0., 3.],
    [1., 1., 0.],
    [0., 2., 0.],
])

q = np.array([1., 0., 1.] )
```

Activity: Document ranking in matrix code (U1-LO4, U1-LO8) (continued)

```
scores = (X @ q) / (  
    np.linalg.norm(X, axis=1) * np.linalg.norm(q)  
)  
ranking = doc_names[np.argsort(scores)[::-1]]  
  
scores, ranking
```

Output:

```
(array([1. , 0.5, 0. ]),  
 array(['D1', 'D2', 'D3']))
```

1. What does the second row of X represent?
2. Which expression computes the three dot products $\mathbf{D}_i \cdot \mathbf{q}$?
3. What does `np.argsort(scores)[::-1]` do?

Example: The attention activity in matrix form

In Gathering context for the bird position (U1-LO1, U1-LO8), we computed the scores

$$\mathbf{q} \cdot \mathbf{k}_{\text{small}}, \quad \mathbf{q} \cdot \mathbf{k}_{\text{red}}, \quad \mathbf{q} \cdot \mathbf{k}_{\text{bird}}$$

one at a time. Matrix-vector multiplication computes the same scores at once.

Put the key vectors as the rows of

$$K = \begin{bmatrix} \mathbf{k}_{\text{small}}^T \\ \mathbf{k}_{\text{red}}^T \\ \mathbf{k}_{\text{bird}}^T \end{bmatrix}.$$

Then

$$K\mathbf{q} = \begin{bmatrix} \mathbf{k}_{\text{small}} \cdot \mathbf{q} \\ \mathbf{k}_{\text{red}} \cdot \mathbf{q} \\ \mathbf{k}_{\text{bird}} \cdot \mathbf{q} \end{bmatrix}.$$

For the numbers in Gathering context for the bird position (U1-LO1, U1-LO8),

$$K = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{bmatrix}, \quad \mathbf{q} = \begin{bmatrix} 1 \\ 1 \end{bmatrix},$$

so

$$K\mathbf{q} = \begin{bmatrix} 1 \\ 1 \\ 2 \end{bmatrix}.$$

Activity: The same attention calculation in code (U1-LO4, U1-LO8)

The following code performs the calculation from the token attention activity.

```
import numpy as np

tokens = np.array(["small", "red", "bird"])

q = np.array([1., 1.])

K = np.array([
    [1., 0.],
    [0., 1.],
    [1., 1.],
])

V = np.array([
    [4., 0.],
    [0., 4.],
    [4., 4.],
])

s = K @ q
alpha = s / s.sum()
out = alpha @ V

s, alpha, out
```

Activity: The same attention calculation in code (U1-LO4, U1-LO8) (continued)

Output:

```
(array([1., 1., 2.]),  
 array([0.25, 0.25, 0.5 ]),  
 array([3., 3.]))
```

1. Which line computes the token scores?
2. Which token receives the largest weight?
3. Which line forms the weighted average?
4. Why do the entries of α add to 1?

Definition: Standard basis

The j -th e_j is the column vector with a 1 in position j and zeros elsewhere.

The term “basis” will be explained later in the course. For now, use the

concrete description of $\mathbf{e}_j$ above.

Fact

If $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$, then $A\mathbf{e}_j = \mathbf{a}_j$. Applying a matrix to a standard basis vector picks out a column.

Why is this true?

The coordinates of $\mathbf{e}_j$ are all zero except for a 1 in position j . Thus the column formula for the matrix-vector product gives

$$A\mathbf{e}_j = \sum_{k=1}^n (\mathbf{e}_j)_k \mathbf{a}_k = \mathbf{a}_j.$$

Here $(\mathbf{e}_j)_k$ denotes the k -th coordinate of $\mathbf{e}_j$, which equals 1 when $k = j$ and 0 otherwise.

Example: The identity matrix leaves vectors unchanged

For every column vector $\mathbf{x} \in \mathbb{R}^n$, $I_n \mathbf{x} = \mathbf{x}$.

Indeed, the columns of the identity matrix from Identity matrix are $\mathbf{e}_1, \dots, \mathbf{e}_n$. Therefore, writing $\mathbf{x} = (x_1, \dots, x_n)$,

$$I_n \mathbf{x} = x_1 \mathbf{e}_1 + \cdots + x_n \mathbf{e}_n = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \mathbf{x}.$$

Geometric matrix actions in $\mathbb{R}^2$: building intuition

In $\mathbb{R}^2$, a matrix $A = [\mathbf{a}_1 \ \mathbf{a}_2]$ sends $\mathbf{e}_1$ to $\mathbf{a}_1$ and $\mathbf{e}_2$ to $\mathbf{a}_2$. Every vector $\mathbf{x} = (x_1, x_2)$ can be written as $x_1 \mathbf{e}_1 + x_2 \mathbf{e}_2$, since

$$x_1 \mathbf{e}_1 + x_2 \mathbf{e}_2 = x_1 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + x_2 \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \mathbf{x}.$$

Therefore,


$$\begin{aligned}
A\mathbf{x} &= A(x_1\mathbf{e}_1 + x_2\mathbf{e}_2) \\
&= x_1A\mathbf{e}_1 + x_2A\mathbf{e}_2 \\
&= x_1\mathbf{a}_1 + x_2\mathbf{a}_2.
\end{aligned}$$

The second equality uses linearity, and the last uses the referenced item. Thus a 2×2 matrix is geometrically determined by where it sends the two coordinate directions.


To understand a 2×2 matrix A , apply it to the four corners of the unit square. The image is usually a parallelogram. The two edges leaving the origin are the columns of A . We will refer to this as

the .


original unit square



columns




image




$$A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$$

Figure. A 2×2 matrix sends the unit square to the parallelogram determined by its two columns. This is the geometric version of columns contribute.


original unit square



columns



image



$$A = \begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix}$$

Figure. A horizontal stretch by a factor of 2 sends the unit square to a rectangle twice as wide. The first column doubles the horizontal basis vector, while the second column leaves the vertical basis vector unchanged.

• **Matrix:** $\begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 2 \\ 0 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Name: horizontal stretch

Unit-square effect: The right edge moves farther right.

• **Matrix:** $\begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Name: projection onto the y-axis

Unit-square effect: Horizontal information is forgotten.

• **Matrix:** $\begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} -1 \\ 0 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Name: reflection across the y -axis

Unit-square effect: The square flips left.

• **Matrix:** $\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

Name: reflection across the line $y = x$

Unit-square effect: The square stays in the same region, but the horizontal and vertical coordinate directions exchange.

• **Matrix:** $\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

Name: horizontal shear

Unit-square effect: The bottom edge stays fixed; the top edge shifts right.

• **Matrix:** $\begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} -1 \\ 0 \end{bmatrix}$$

Name: 90° counterclockwise rotation

Unit-square effect: The square rotates around the origin.

• **Matrix:** $\begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$

$$Ae_1: \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$Ae_2: \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Name: projection onto the x -axis

Unit-square effect: Vertical information is forgotten.

Practice recognizing such geometric transformations in Shear Madness.

Activity: Horizontal shear (U1-LO4)

Let $A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$. Compute Ae_1 , Ae_2 , and Ax , where $x = (1, 1)$.

Note

For the shear matrix $A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$, the row view gives

$$A \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} (1, 1) \cdot (x, y) \\ (0, 1) \cdot (x, y) \end{bmatrix} = \begin{bmatrix} x + y \\ y \end{bmatrix}.$$

The first row measures horizontal plus vertical, while the second row measures vertical only. The column view gives


$$A \begin{bmatrix} x \\ y \end{bmatrix} = x \begin{bmatrix} 1 \\ 0 \end{bmatrix} + y \begin{bmatrix} 1 \\ 1 \end{bmatrix},$$

so the y -coordinate contributes both upward and rightward motion.


Activity: Reflection across the y -axis (U1-LO4)

Let $A = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$. Compute $A\mathbf{e}_1$, $A\mathbf{e}_2$, $A\mathbf{x}$, and $A\mathbf{y}$, where $\mathbf{x} = (1, 1)$ and $\mathbf{y} = (0, 1)$.


original unit square



columns



image



$$A = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$

Figure. Reflection across the y -axis sends the unit square to its left-hand mirror image. The first column reverses the horizontal direction, while the second column leaves the vertical direction unchanged.

Activity: Reflection across the line $y=x$ (U1-L04, U1-L05)


Let

$$A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}.$$


1. Compute Ae_1 and Ae_2 .
2. Compute $A \begin{bmatrix} x \\ y \end{bmatrix}$.
3. What happens to a point on the line $y = x$?




original unit square



columns



image



$$A = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$$

Figure. Three-panel unit-square visualization of reflection across y equals x , with exchanged basis arrows.

Linear and affine maps

A function assigns one output to each input. In this course, many functions have vectors as inputs and vectors as outputs:

$$F : \mathbb{R}^n \rightarrow \mathbb{R}^m.$$

The of F is its set of allowed inputs, $\mathbb{R}^n$, and its is the specified set containing its outputs, $\mathbb{R}^m$. A matrix gives one important source of such functions by the rule $\mathbf{x} \mapsto A\mathbf{x}$.

Definition: Matrix map

Given an $m \times n$ matrix A , the induced by A is the function $T_A : \mathbb{R}^n \rightarrow \mathbb{R}^m$ defined by

$$T_A(\mathbf{x}) = A\mathbf{x}.$$

Definition: Linear map

A function $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is if

$$\begin{aligned} T(\mathbf{u} + \mathbf{v}) &= T(\mathbf{u}) + T(\mathbf{v}), \\ T(c\mathbf{v}) &= cT(\mathbf{v}) \end{aligned}$$

for all vectors $\mathbf{u}, \mathbf{v}$ and all scalars c .

Note: Linear maps preserve linear combinations

If $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is linear, then for any vectors $\mathbf{v}_1, \dots, \mathbf{v}_k \in \mathbb{R}^n$ and scalars $c_1, \dots, c_k \in \mathbb{R}$,

$$T(c_1\mathbf{v}_1 + \dots + c_k\mathbf{v}_k) = c_1T(\mathbf{v}_1) + \dots + c_kT(\mathbf{v}_k).$$

Thus a linear combination of inputs is sent to the linear combination of their outputs with the same coefficients. This follows by repeatedly applying the two linearity rules.

Theorem: Matrix maps are linear

If A is an $m \times n$ matrix, then the function

$$T_A : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad T_A(\mathbf{x}) = A\mathbf{x}$$

is linear.

Why is this true?

Matrix-vector multiplication distributes over vector addition and scalar multiplication:

$$A(\mathbf{u} + \mathbf{v}) = A\mathbf{u} + A\mathbf{v}, \quad A(c\mathbf{v}) = cA\mathbf{v}.$$

These are exactly the two linearity rules.

Every vector $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$ can be written as a linear combination of the standard basis vectors. Indeed, $x_j \mathbf{e}_j$ has entry x_j in position j and zeros elsewhere, so adding these vectors coordinate by coordinate gives

$$x_1 \mathbf{e}_1 + \dots + x_n \mathbf{e}_n = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \mathbf{x}.$$

Since a linear map preserves linear combinations, its values on $\mathbf{e}_1, \dots, \mathbf{e}_n$ determine its value on every input. This observation leads to the following theorem.

Theorem: Every linear map has a matrix

Every linear map $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is given by multiplication by a unique matrix. The columns of that matrix are $T(\mathbf{e}_1), \dots, T(\mathbf{e}_n)$.

Why is this true?

Using the standard-basis expansion above, let $A = [T(\mathbf{e}_1) \ \dots \ T(\mathbf{e}_n)]$. By linearity and the column view of matrix-vector multiplication,

$$\begin{aligned} T(\mathbf{x}) &= T(x_1 \mathbf{e}_1 + \dots + x_n \mathbf{e}_n) \\ &= x_1 T(\mathbf{e}_1) + \dots + x_n T(\mathbf{e}_n) \\ &= [T(\mathbf{e}_1) \ \dots \ T(\mathbf{e}_n)] \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \\ &= A\mathbf{x}. \end{aligned}$$

The second equality uses linearity, and the third uses the column view of matrix-vector multiplication. This matrix is unique because its j -th column must be $A\mathbf{e}_j = T(\mathbf{e}_j)$.

Activity: Images of basis vectors determine the matrix (U1-L05)

Suppose a linear map $T : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ satisfies $T(\mathbf{e}_1) = (2, 1)$ and $T(\mathbf{e}_2) = (-1, 3)$. Find the matrix A such that $T(\mathbf{x}) = A\mathbf{x}$.

Warning

We use “linear map” and “linear transformation” interchangeably; “transformation” is often used when emphasizing geometry.

Definition: Affine map

Let A be an $m \times n$ real matrix and let $\mathbf{b} \in \mathbb{R}^m$ be a fixed column vector. The function

$$F : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad F(\mathbf{x}) = A\mathbf{x} + \mathbf{b},$$

where $\mathbf{x} \in \mathbb{R}^n$ is a column vector, is called . It is built from

a linear map followed by a .

An affine map is linear only when $\mathbf{b} = \mathbf{0}$. A quick test is the zero vector: every linear map sends $\mathbf{0}$ to $\mathbf{0}$, but

$$A\mathbf{0} + \mathbf{b} = \mathbf{b}.$$

If $\mathbf{b} \neq \mathbf{0}$, the map is not linear.




Figure. A unit square is stretched horizontally and then translated one unit right and one unit up.

Example: An affine layer

A layer in a neural network can take the form

$$\mathbf{x} \mapsto W\mathbf{x} + \mathbf{b}.$$

The symbol $\mapsto$ means “maps to”: the input vector $\mathbf{x}$ is sent to the output vector $W\mathbf{x} + \mathbf{b}$. The matrix W mixes the input coordinates. The vector $\mathbf{b}$ shifts the result. Since

$$W\mathbf{0} + \mathbf{b} = \mathbf{b},$$

this map is affine. It is linear exactly when $\mathbf{b} = \mathbf{0}$.